



```
In [1]: import nltk
        from nltk import word_tokenize
        import re
        from nltk.corpus import stopwords
        import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        from nltk.stem import PorterStemmer, LancasterStemmer
        from nltk.stem import WordNetLemmatizer
        from sklearn.feature_extraction.text import CountVectorizer
        from nltk import word_tokenize, sent_tokenize
        from sklearn.linear_model import LogisticRegression
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.naive_bayes import MultinomialNB
```

Read the Dataset

```
In [2]: df=pd.read_csv('Downloads\\amazon_alex.csv')
```

```
In [3]: df.head()
```

Out[3]:

	reviews	sentiment
0	Love my Echol	1
1	Loved it!	1
2	Sometimes while playing a game, you can answer...	1
3	I have had a lot of fun with this thing. My 4 ...	1
4	Music	1

Handling Null Values(If any)

```
In [4]: df.isnull().sum()
```

```
Out[4]: reviews      0  
sentiment      0  
dtype: int64
```

Preprocess the Data

```
In [5]: sw = stopwords.words('English')  
lm = WordNetLemmatizer()
```

```
In [6]: data = []
for i in df['reviews']:
    t = i.lower()           # lower case conversion
    t = re.sub('[^A-Za-z]', ' ', t)  # removing punctuation
    t = word_tokenize(t)     # word tokenization
    t = [i for i in t if i not in sw]  # stop words removal
    t = [lm.lemmatize(i, pos="v") for i in t]  # lemmatization # returns list of words
    t = " ".join(t)         # Joining all the words
    data.append(t)
```

Transforming the Words into Vectors

```
In [7]: cv=CountVectorizer()
sm=cv.fit_transform(data).toarray()
print(sm)
print(cv.get_feature_names_out())
print(len(cv.get_feature_names_out()))

[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 ...
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]
['abay' 'abc' 'abd' ... 'zonked' 'zzzz' 'zzzzzzz']
3086
```

```
In [8]: df2=pd.DataFrame(sm,columns=cv.get_feature_names_out())
df2.head()
```

```
Out[8]:
```

	abay	abc	abd	abide	abilities	ability	able	absolutely	absolutly	ac	...	youngest	youtube	yr	yrs	yup	zero	zigbee	zonked	zzzz	zzzzzzzz
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	1	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

5 rows x 3086 columns

```
In [9]: df3 = pd.DataFrame (data, columns = ['reviews'])
df['reviews']=df3['reviews']
df
```

```
Out[9]:
```

	reviews	sentiment
0	love echo	1
1	love	1
2	sometimes play game answer question correctly ...	1
3	lot fun thing yr old learn dinosaurs control l...	1
4	music	1

3145	perfect kid adults everyone	1
3146	listen music search locations check time look ...	1
3147	love things run entire home tv light thermosta...	1
3148	complaint sound quality great mostly use comma...	1
3149	good	1

3150 rows x 2 columns

```
In [10]: x = cv.fit_transform(data).toarray()
y = df.iloc[:, 1].values
```

Splitting the data into Training and Testing Data

```
In [11]: x_train,x_test,y_train,y_test = train_test_split(x,y)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```


```
(2362, 3086)
(788, 3086)
(2362,)
(788,)
```



Applying Multinomial Naive Bayes Classification

```
In [12]: m1=MultinomialNB();  
m1.fit(x_train,y_train)
```

```
Out[12]: MultinomialNB()
```




```
In [14]: print('Training score=',m1.score(x_train,y_train))  
         print('Testing score=',m1.score(x_test,y_test))
```

```
Training score= 0.958086367485182  
Testing score= 0.9378172588832487
```

```
In [15]: print('accuracy score=',accuracy_score(y_test,ypred_m1))
```

```
accuracy score= 0.9378172588832487
```

Applying Logistic Regression

```
In [16]: m2=LogisticRegression();  
         m2.fit(x_train,y_train)
```

```
Out[16]: LogisticRegression()
```



```
In [17]: ypred_m2=m2.predict(x_test)
print(ypred_m2)
```

[illegible]

```
In [18]: print('training score',m2.score(x_train,y_train))  
         print('testing score',m2.score(x_test,y_test))
```

```
training score 0.9767146486028789  
testing score 0.9467005076142132
```

```
In [19]: print('accuracy score=',accuracy_score(y_test,ypred_m2))
```

```
accuracy score= 0.9467005076142132
```

Applying KNN Model

```
In [20]: m3=KNeighborsClassifier(n_neighbors=12)  
         m3.fit(x_train,y_train)
```

```
Out[20]: KNeighborsClassifier(n_neighbors=12)
```

```
In [21]: ypred_m3=m3.predict(x_test)
print(ypred_m3)
```

[illegible]

```
C:\Users\dell\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
```

C:\Users\dell\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
In [22]: print('Training score=', m3.score(x_train, y_train))
print('Testing score=', m3.score(x_test, y_test))
```

C:\Users\dell\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
Training score= 0.9123624047417442
```

```
Testing score= 0.9200507614213198
```

C:\Users\dell\anaconda3\lib\site-packages\sklearn\neighbors_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
In [23]: print('accuracy score=', accuracy_score(y_test, ypred_m3))
```

```
accuracy score= 0.9200507614213198
```

Computing Confusion Matrix and Classification Report

Multinomial Naive Bayes model

```
In [24]: cm1=confusion_matrix(y_test,ypred_m1)
print(cm1)
print(classification_report(y_test,ypred_m1))
```

```
[[ 16  43]
 [  6 723]]
```

	precision	recall	f1-score	support
0	0.73	0.27	0.40	59
1	0.94	0.99	0.97	729
accuracy			0.94	788
macro avg	0.84	0.63	0.68	788
weighted avg	0.93	0.94	0.92	788

Logistic Regression

```
In [25]: cm2=confusion_matrix(y_test,ypred_m2)
print(cm2)
print(classification_report(y_test,ypred_m2))
```

```
[[ 21  38]
 [  4 725]]
```

	precision	recall	f1-score	support
0	0.84	0.36	0.50	59
1	0.95	0.99	0.97	729
accuracy			0.95	788
macro avg	0.90	0.68	0.74	788
weighted avg	0.94	0.95	0.94	788

KNN Model

```
In [26]: cm3=confusion_matrix(y_test,ypred_m3)
print(cm3)
print(classification_report(y_test,ypred_m3))
```

```
[[ 0 59]
 [ 4 725]]
```

		precision	recall	f1-score	support
	0	0.00	0.00	0.00	59
	1	0.92	0.99	0.96	729
accuracy				0.92	788
macro avg		0.46	0.50	0.48	788
weighted avg		0.86	0.92	0.89	788

Best Accuracy Model

From the above Models, the Logistic Regression Model has the best accuracy